UDAB-ViS: User Driven Adaptable Bandwidth Video System

Dustin Wright · Yusuf Ozturk

Abstract Adaptive bitrate (ABR) streaming has become an important and prevalent feature in many multimedia delivery systems, with content providers such as Netflix and Amazon using ABR streaming to increase bandwidth efficiency and provide the maximum user experience when channel conditions are not ideal. Where such systems could see improvement is in the delivery of live video with a closed loop cognitive control of video encoding. In this research, we present a camera system which provides spatially and temporally adaptive video streams, learning the users preferences in order to make intelligent scaling decisions. The system employs a hardware H.264/AVC encoder for video compression. The encoding parameters can be configured by the user or by the cognitive system on behalf of the user when the bandwidth changes. The cognitive video client developed in this study learns the users preferences (i.e. video size over frame rate) over time and intelligently adapts encoding parameters when the channel conditions change. It has been demonstrated that the cognitive decision system developed has the ability to control video bandwidth by altering the spatial and temporal resolution, as well as the ability to make scaling decisions.

D. Wright
San Diego State University, Electrical and Computer Engineering Department, 5500 Campanile Dr. San Diego, California 92182
E-mail: wright21@rohan.sdsu.edu

Y. Ozturk
San Diego State University, Electrical and Computer Engineering Department, 5500 Campanile Dr. San Diego, California 92182

1 Introduction

As wireless networks become more ubiquitous and the number of devices capable of accessing these networks increases, the need for more efficient video streaming solutions becomes vastly important. By 2015, it is expected that approximately 90 percent of online consumer traffic and almost 66 percent of mobile traffic will be video [8]. With an increasing amount of video traffic, bandwidth efficiency becomes a serious concern in order to deliver the best quality of experience (QoE) to each individual user. At the same time, servers should be able to deliver video in such a way that information the user deems important is not lost due to bandwidth constraints and the method by which the video adapts.

The rest of the paper is structured as follows; in section 3 we review H.264 packetization. In section 3.2, we review H.264/SVC as a scalable streaming solution, as well as the scalability solution used in our system. Sections 5 and ?? will describe the system architecture proposed in this study, as well as the method by which the video is encoded. Section 4 will detail our learning model and how video bandwidth is optimized. Finally, test setup and experimental results will be presented in section 6, and we conclude in section 7.

2 Related Works

The problem of sending a continous video stream over an uncertain channel is not new and a range of solutions have been proposed. Many of the approaches are based on already existing protocols such as the Real-Time Streaming Protocol RTSP [11], TCP, and HTTP [8,?,?,?] and achieve bandwidth adaptability in one of a few ways. These bandwidth adaptability techniques include progressive download, adaptive bitrate (ABR) streaming, and stream-switching [4]. In progressive download, video is transmitted as regular data files using TCP and is buffered by the client; playing starts when a sufficient amount of buffering has been achieved. With ABR, the server selects the encoding bitrate in order to optimize the video's SNR resolution for a given channel. The result is a drastic reduction in the need for buffering; from the end user's perspective, the quality resolution of the video changes as network conditions change. Examples of ABR solutions are found in HTTP Adaptive Streaming (HAS) [8] and Dynamic Adaptive Streaming over HTTP (DASH) [6], used in the popular video content provider Netflix [2]. Finally, with stream switching, the server encodes the source video with different encoding parameters and allows the client to switch between streams based on network conditions. Examples of stream switching solutions can be found in [4],[1] and [18]. Such techniques are ideal for online video streaming as they can use HTTP to negotiate streaming parameters and transmit the video stream; however, the major pitfall is that in most cases, only the video bitrate will be affected and no control is exercised over the spatial and temporal resolution of the video. In the case where spatial and temporal resolution can be affected, raw video

will have to be re-encoded or transcoded at the source which can cause a delay in the video being transmitted [4].

In addition to protocol based approaches which tend to only allow for bitrate scalability, many codec based approaches have been developed which allow for easy spatial and temporal scalability [7]. Two prime examples of codec based approaches are H.264/AVC and H.264/SVC [12][14]. With these codecs, video need only be encoded once and, due to the nature of the decoders, can be transmitted as several sub-streams in order to control the temporal or spatial resolution. The primary issue with this approach is that it limits the client to only a certain set of video decoders. A review of the H.264/SVC approach to scalability will be presented later in this paper.

A recent area of interest seen in the literature is the autmoation of video encoding parameter selection as a means to provide both exceptional quality of service (QoS) and quality of experience (QoE). Approaches to this problem include both cognitive and non-cognitive solutions, though fewer cognitive systems have yet been proposed. In [4], the authors propose a *Quality Adaptation Controller* which uses a proportional-integral (PI) controller at the server to select an appropriate video stream for the client. The system uses stream switching and employs feedback control to maintain the video bandwidth below the available channel bandwidth. Examples of cognitive solutions are found in [10], which use statistical models that adapt to user feedback in order to make encoding parameter selections on their behalf.

This study investigates a cognitive approach to video bandwidth control based on user preferences learned by C support vector machines (SVM). This has an advantage over previous cognitive solutions in that it takes into account multiple features related to the video, such as video content, and is scalable in that more features may be observed if needed. We present a solution that is novel in its application of machine learning as an accurate method to learn user preferences.

3 H.264 Relevant Background and Proposed Scaling Architecture

3.1 H.264/AVC Basics

H.264/AVC is a video coding standard developed jointly by the ITU-T Video Coding Experts Group (VCEG) and ISO/IEC Moving Pictures Expert Group (MPEG) designed with the goals of enhanced video compression and "network friendly" video representation addressing "conversational" applications such as video conferencing, as well as "non-conversational" applications such as broadcast streaming [17]. In December of 2001 VCEG and MPEG formed a Joint Video Team (JVT) which in March of 2003 finalized the draft of the H.264/AVC video coding standard for formal submission [17].

The standard provides highly efficient video coding andis used in a breadth of applications, from storage to streaming. It provides bitrate savings of 50% or more over its predecessor video codecs [16], making H.264/AVC especially

applicable in wireless environments. The effectiveness of H.264/AVC as a tool
for video compression over IP networks and in wireless settings is reviewed
in [15] and [13]. In addition, H.264 employs a litany of features to enhance
the quality of video coding and error resiliency over previous standards. [17]
and [5] provide a detailed overview of these features. The basic structure of
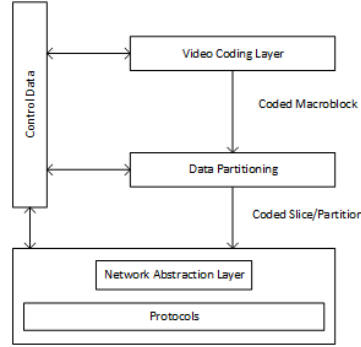an H.264/AVC encoder is depicted in Figure 1.

Fig. 1 H.264/AVC Encoder Structure [17]

The codec given in Figure 1 covers both the Video Coding Layer (VCL)
and the Network Abstraction Layer (NAL). The VCL performs the physical
encoding and compression of video while the NAL wraps a header around video
packet data in order to assist the decoder in understanding how to handle the
packetized frames. We will now discuss the NAL and give a general overview
of how frames are packetized when sent using UDP/RTP.

The NAL allows the ability to map and packetize data with a multitude
of transport layer protocols (i.e. UDP/RTP, file formats, etc.). When frames
are encoded in the VCL they are organized into NAL units which serve as a
wrapper to the underlying information. Each NAL unit contains a header byte
that indicates what type of data is contained in this unit. This allows for the
segmentation of video into packets, with the NAL unit indicating the start of
a new access unit. The NAL unit contains a one byte header and a payload
byte string [16]. The header indicates the type of NAL unit, potential presence
of errors, and information about the relative importance of this NAL unit in
the decoding process [16]. The structure of the one-byte NAL unit header is
shown in Figure 2.

NAL Unit Header

F (1 bit)   NRI (2 bits)   Type (5 bits)

Fig. 2 NAL Unit Header Structure

The fields of the header are designated as follows:

– F: forbidden bit; should always be 0
– NRI: used to indicate if the content of this NAL unit should be used to reconstruct reference pictures in inter picture prediction
– Type: Specifies the NAL unit payload type

Examples of NAL unit types are parameter sets that contain relevant information about a the video stream or an individual frame, as well as frame type (I frame, P frame, B frame) and slice priority.

The H.264/AVC specification also defines a set of profiles and levels which specify different sets of required functional support for decoders. According to [17], "A profile defines a set of coding tools or algorithms that can be used in generating a conforming bit-stream, whereas a level places constraints on certain key parameters of the bitstream." For video conferencing applications or streaming from mobile devices such as the system proposed in this paper, the constrained baseline or baseline profile are appropriate choices. More detailed information on profile types and their constraints can be found in section A.2 of [5].

In the next section, we comment on the scalable video coding extension of H.264 and compare it with the cognitive parameter adaptation method propsed in this paper.

3.2 Video Scaling Architecture

The need for scalable video codecs can be characterized by the following scenario: when transmitting a stream at a certain quality with a video bandwidth B over a congested channel where the channel bandwidth C fluctuates such that $C < B$, the receiving terminal may experience significant degradation of video quality. Scalable codecs combat this by altering one or more resolutions of the video in order to fit the channel. H.264/SVC is once such codec that uses composable bit streams as a means to scale video. Certain parts of the bit stream are removed, separating one stream into layered substreams in such a way that the underlying streams are still decodable [12]. For example, a transmitter may send one base layer bit stream and multiple enhancement layer bit streams with the receiver selecting which of these to send to the decoder. In this, video bandwidth can be controlled by choosing only the necessary bit streams to stay within the channel bandwidth.

SVC presents a sharp contrast to classic single layer video streams in which one decodable bit stream is transmitted. In order to have control over the bandwidth of the video, the transmitter must encode the source video with different encoding parameters and the receiving decoder must adapt to these changes. We will next summarize the Scalable Video Coding extension of H.264 as described in [12] and compare it to the scaling method developed in this study that used a single layer video stream, altering encoding parameters at the source.

The Scalable Video Coding extension of H.264/AVC inherits all of the base functionality of H.264 with only the necessary added features to achieve scalable video streaming. In this, it supports the primary scalability parameters, being temporal, spatial, and quality resolution. To achieve temporal scalability, the transmitter may send multiple temporal streams divided into a temporal base layer and one or more temporal enhancement layers [12]. One may label these streams as T_0 through T_k. A receiving decoder then simply needs to know which of these access units are valid or invalid for the current stream, starting from 0 through n where $n \leq k$. The ability to partition a stream as such and play only the valid streams is already present in the H.264/AVC standard with the employment of reference picture memory control [12]. The partitioning of a stream into multiple temporal streams is illustrated in Figure 3.


Fig. 3 Temporal Scalability with H.264/SVC [12]

In order to achieve spatial scalability, SVC uses multi-layer coding with inter-layer prediction [12]. Multiple layers are transmitted, each corresponding to a specific spatial resolution and referred to by an integer valued dependency identifier between 0 and $d - 1$ where d is the number of spatial layers [12]. Quality scalability works on the same principle as spatial scalability with the layers transmitted being of the same spatial resolution.

In the proposed single-layer coded video stream, one bit stream is encoded and sent to the receiver. This bit stream is of a fixed spatial, temporal, and quality resolution for the entire sequence of video. There is no mechanism inherent to the codec to change the spatial, temporal, or quality resolution midstream. In this, the video bandwidth is controlled by switching encoding parameters at the source, effectively segmenting the video in time. The video stream is partitioned into multiple sequences labeled $T_k, k = 0 \ldots n - 1$ where n is the number of segments for the given session and T_k is the time instance when the segment k begins. The partitioning is determined on the fly as a

function of the channel bandwidth, where channel bandwidth can be measured with a reasonable degree of accuracy; for example, using a method like DIChirp as laid out in [9]. At each time instance T_k, the transmitter resets the encoding parameters in such a way that the bandwidth of the video is altered to fit to the channel. When this occurs, a new sequence parameter set is introduced into the stream to signal to the decoder the changes to the encoded video. A sample stream is depicted in Figure 4. The changes to the
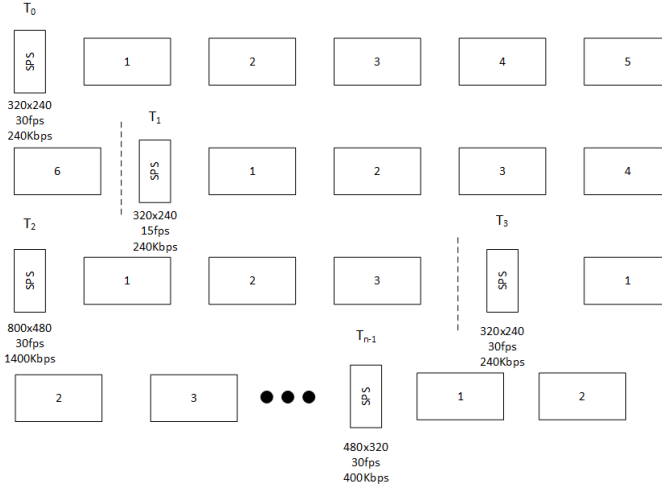

Fig. 4 Single Layer Encoder Parameter Switching

encoding parameters will insert a delay into the stream for the time it takes to restart the encoding process. To compensate for this delay and to handle the alteration events at time T_k, we propose the receiver architecture in Figure 5. In the proposed architecture, a preprocessor inspects the NAL unit headers of


Fig. 5 Receiver Architecture

each incoming packet, waiting for a new sequence parameter set. At this event, the idle decoder thread is invoked and set up to decode the next sequence of video. The previously active decoder thread empties its queue prior to the start of the new segment of video. When the active thread signals completion, the new thread takes over. This effectively mitigates any delay that may be

introduced due to reconfiguring the decoder, providing a smooth stream for the user.

Delay on the encoder side is greatly avoided by using a hardware H.264/AVC encoder. Encoder initialization happens in real time and there is no CPU overhead when it comes to encoding frames. In addition, this real-time efficiency allows for minimal delay between video segments. Finally, our method allows us to control the temporal, spatial, and quality resolution of each video segment in a more granular fashion than SVC as we are not restricted to a discrete number of video layers.

4 User Profiles and Bandwidth Optimization

The development of user profiles allows the client to make intelligent scaling decisions in line with how the user would have changed the video. This relieves the burden on the user to figure out how video should be scaled in different contexts, making it critical for a widely acceptable system. The learning and prediction mechanism proposed in this study offers a simple, yet effective way to provide dynamically adaptive video streams tailored to each individual.

4.1 Preferences and Profiles

A profile of the user is developed by the client application which will determine how video is scaled when it is no longer optimized to the transmission channel. In this, a developed profile places the user in one of four discrete classes which represent their preferences in relation to video quality:

- Class 0: User prefers high temporal and spatial resolution, no quality preference.
- Class 1: User prefers high temporal resolution; optimize spatial resolution for quality.
- Class 2: User prefers high spatial resolution; optimize temporal resolution for quality.
- Class 3: User prefers optimal quality; configure temporal and spatial resolution appropriately.

Knowing these preferences, the video bandwidth is optimized by weighing the video resolutions with higher or lower priority. This translates into the client treating lower priority parameters more harshly when determining the new encoding parameters.

Transforming these classes to their equivalent binary form we can represent them with a 2 bit value in which the first bit conveys the temporal resolution preference and the second bit conveys the spatial resolution preference. A set bit indicates the desire for a high weight given to this resolution at the expense of quality, and an unset bit indicates the desire to optimize quality resolution at the expense of the resolution in question. In this, we can find the best way to

alter the spatial and temporal resolution such that the video bandwidth will fit the channel and align with the user's desires in terms of quality. The creation of these user profiles is accomplished using machine learning; in particular, by solving the classification problem using support vector machines.

4.2 Creating User Profiles

In order to create a profile for each user, we employed a supervised learning algorithm to create a decision function based on the users behavior to predict the class a user falls into. The decision function is calculated using the LibSVM implementation of a C support vector machine with a radial basis kernel [3].

In supervised machine learning, a training set composed of a series of training samples is presented as input to the learning algorithm, the output of which is a set of coefficients for a decision function. In classification problems, the training samples are the values of a set of features which are relevant to the output, and a label for each of these vectors which denotes what class applies at the instant of the sample. The features used in this study are the channel bandwidth and the content type of the video (i.e. high quality medical, talking head, sporting event, etc.). Two C support vector machines with radial basis kernels are used to learn the two preferences. The C support vector machine classification algorithm used in this study is defined in [3]. To use support vector machines to learn user preferences, a training set is defined as a feature vector $x_i \in \mathbf{R^n}, i = 1 \ldots l$ and the class label vector $y_i \in \{1, -1\}$ where 1 and -1 indicate the two distinct classes. The primal optimization problem in equation (1) is then solved [3]

$$\begin{aligned} &\underset{w,b,\xi}{\text{minimize}} \quad \frac{1}{2}w^T w + C\sum_{i=1}^{l}\xi_i \\ &\text{subject to} \quad y_i(w^T\phi(x_i) + b) \geq 1 - \xi_i \\ &\qquad\qquad \xi_i \geq 0, i = 1, \ldots, l \end{aligned} \tag{1}$$

where $\phi(x_i)$ maps x_i into higher dimensional space and $C > 0$ is a configurable parameter [3]. The dual problem presented in equation (2) can then be solved in order to account for the possible high dimensionality of the vector parameter w.

$$\begin{aligned} &\underset{\alpha}{\text{minimize}} \quad \frac{1}{2}\alpha^T Q\alpha - e^T\alpha \\ &\text{subject to} \quad y^T\alpha = 0 \\ &\qquad\qquad 0 \leq \alpha_i \leq C, i = 1, \ldots, l \end{aligned} \tag{2}$$

where e is a column vector of all ones, $Q_{(i,j)} = y_i y_j K(x_i, x_j)$ and $K(x_i, x_j) \equiv \phi(x_i)^T\phi(x_j)$ is the radial basis function (RBF). The RBF K is a Gaussian distribution, presented in equation (3).

$$K(x_i, x_j) = e^{-\gamma||x_i - x_j||^2} \tag{3}$$

where γ is a configurable parameter selected by the user. Finally, the optimal w is computed using equation (4), and the decision function is laid out in equation (5).

$$w = \sum_{i=1}^{l} y_i \alpha_i \phi(x_i) \qquad (4)$$

$$sgn(w^T \phi(x) + b) = sgn(\sum_{i=1}^{l} y_i \alpha_i K(x_i, x_j) + b) \qquad (5)$$

Considering the graphical representation of x_i and y_i, we have a multidimensional input space and class labels associated with each vector from x. The support vector machine attempts to find a hyperplane that separates the two classes with the widest margin, using kernels to support nonlinear separations. From this hyperplane comes the decision function in equation (5) that is used to predict class labels for new input vectors. In this study, we train the support vector machines to predict the users preferences using a combination of implicit and explicit feedback.

When a new user begins interacting with the system, they are initiated in "learning mode." In learning mode, a change in channel bandwidth will result in a knee-jerk reaction by the system to simply alter the quality resolution of the video, leaving spatial and temporal resolution unchanged. Explicit feedback is obtained from the user. The user's interaction with the system is recorded and forms the training data set for the support vector machines. When the training set is sufficiently large enough to accurately train the support vector machines, 3-fold cross validation is performed to ensure accurate selection of C in equation (1) and γ in equation (3) [3]. From this, a decision function is created which is used to make predictions. The explicit feedback process is shown in Figure 6.
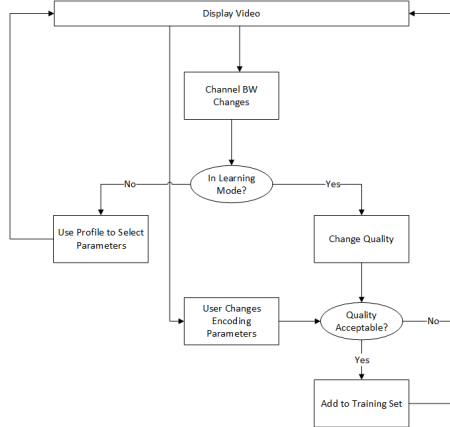

Fig. 6 Feedback Process

4.3 Video Bandwidth Determinations

The catalyst for video bandwidth recalculation is when the channel conditions have changed significantly enough to warrant scaling the video. Several methods, such as DIChirp [9], have been shown to estimate the bandwidth of the channel within a high degree of accuracy. We will now lay out how changing the various video resolutions will affect the video bandwidth in our system.

The compression bitrate of the video changes in order to change the quality of the video. The maximum bitrate can simply be calculated as a function of the known channel bandwidth and an optimization constant:

$$Bitrate_{max} = Bandwidth_{channel} * K \tag{6}$$

where $Bitrate_{max}$ is the maximum allowable bitrate, $Bandwidth_{channel}$ is the bandwidth of the channel, and K is the optimization constant determining the percentage of available bandwidth that is acceptable to fill. Due to limitations in our encoder, the max bitrate must be calculated as a function of the frame rate. The following calculation is used:

$$Bitrate_{max} = \frac{Bandwidth_{channel}}{\frac{T}{17.28}} * K \tag{7}$$

where T is the frame rate of the video. The optimal bitrate for a given spatial resolution is then found as a function of the video dimensions and an optimization constant:

$$Bitrate_{opt} = width * height * Q_{max} \tag{8}$$

where $width$ and $height$ are the dimensions of the video, and Q_{max} is the optimization constant which can be configured to find the highest necessary bitrate to deliver high quality video at a given spatial resolution. The output bitrate is then simply $Bitrate_{max}$ unless we wish to be conservative with channel bandwidth. In this case, the best choice of bitrate is either the maximum allowable bitrate or maximum necessary bitrate for the given channel, spatial resolution, and temporal resolution. The output bitrate is then selected from either equation (9) or (10).

$$Bitrate_{out} = Bitrate_{max} \tag{9}$$

$$Bitrate_{out} = \min(Bitrate_{max}, Bitrate_{opt}) \tag{10}$$

The highest necessary bitrate for delivering the best possible quality to the client is given in equation (10).

Determining the temporal and spatial resolution will affect the quality resolution such that higher values will result in lower quality once the bitrate has reached $Bitrate_{max}$. The user's class is inspected in order to set these two parameters, as the class denotes if the temporal/spatial resolution should be optimized or if the quality should be optimized. In this study, we limit the possible temporal resolution values to 30 and 15. In the case where the

temporal resolution bit is set in the users class, the frame rate simply becomes 30 fps, and if unset, the frame rate becomes 15 fps. If the spatial resolution bit is set, the spatial resolution is changed to the maximum value. When the bit is unset, the system first calculates $Bitrate_{max}$, then reduces the spatial resolution until:

$$Bitrate_{opt} \leq Bitrate_{max} \tag{11}$$

This results in a spatial resolution that will be small enough to provide optimal quality video.

4.4 Scaling Decisions

The way that video bandwidth is optimized depends on the class that the user falls into for the given feature set at the time T_k (the moment when the encoding parameters of the k^{th} segment of video are selected). When channel bandwidth changes significantly enough, the first action taken is to predict the users preferences using the two decision functions already generated. This yields a bitmask that serves as the users class; the primary purpose of this class is to specify the order in which to scale the encoding parameters. This ordering also depends on whether or not the video bandwidth is increasing or decreasing. The chart given in Figure 7 depicts the possible decisions that can be made about the video bandwidth.
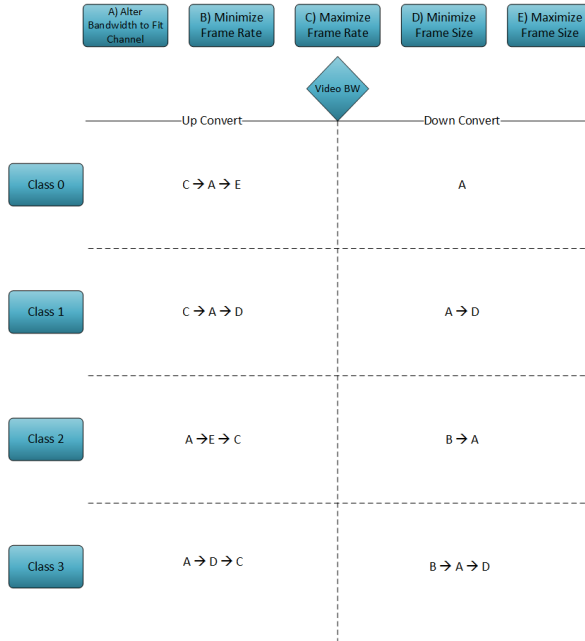

Fig. 7 Bandwidth Decisions

The client application can then successfully adapt to channel bandwidth changes and alter the encoding parameters in such a way that the users desires are fulfilled.

5 System Architecture

Fig. 8 System Architecture

A basic outline of our system architecture is depicted in Figure 8, with the streaming clients, streaming server, and constituent components. The camera system developed in this study consists of a video client and a streaming server. The client connects to the server to request a new streaming session, at which time live video transmission begins. The streaming server is an Android application implemented on Qualcomm MSM8960 hardware. Video is encoded on the device using a hardware H.264/AVC encoder and streamed to the client using raw UDP packets. The client and server designs are detailed in the following sections.

5.1 Client Design

The client uses LibVLC, a library used in the popular VLC media player developed by the VideoLAN group. LibVLC is equipped with all major functionality required to play our live video streams and handles the decoding of H.264 frames. Through configuration of the media player back end, we set up two decoder threads that are used to demultiplex an H.264 stream, consistent with the architecture described in section **??**. This enables the dynamic change of encoding parameters midstream without seeing a significant effect on video playback.

On top of the lower level streaming layer is a signaling layer utilizing TCP for communication between the server and client. When a new video stream is requested, the client attempts to this TCP connection with the server, and upon successful connection, starts the media player. This client then follows a simple protocol to interact with the server which will be discussed in section 5.3.

The client is responsible for determining the correct choice of encoding parameters and for managing the bandwidth of the video based on the current

network capacity. Encoding parameter decisions are based on the user's preferences using the proposed learning model. If an accurate model of the users preferences has been developed, this model drives the parameter selection. A default decision function is employed when learning the users preferences. The client is also intelligent enough not to interfere with the user when they make their own decisions about how to scale the video.

5.2 Server Design

Our camera server application captures live video from an 8MP camera, at varying spatial and temporal resolutions. A TCP server handles all incoming connections from clients and services any requests. On each connection request, a new thread is forked, creating an interface between the client and server. A handle to the encoder is given to each of these threads to allow them control over the resolutions of the video streams. The handle is encapsulated in an object we call the "encoder activation interface". This object, as the name implies, acts as an interface to the encoder (as well as the camera). Via this interface a consumer may initialize, destroy, and alter an encoder for a certain video stream. This allows the clients full control over the parameters of the video, including the video bandwidth. The server remains agnostic of channel conditions and acts as a slave to the connected clients, reconfiguring the stream as necessary based on the request, because the client application learns the users preferences and therefore makes more intelligent decisions about the encoding parameters.

5.3 Session Management

The system contains a signaling layer using TCP for messaging between the client and server. This layer acts as a session manager. TCP is used for reliable communication of messages between terminals as well as to signal the beginning and end of a streaming session. A streaming session begins once the server accepts a client's connection, and ends when one of the terminals disconnects. When a client wishes to receive a particular stream from a server, it first attempts to make a connection with the server. Upon connection, the server initializes a new thread to service the client's requests. The server thread first starts the encoder, which begins streaming packets containing the encoded H.264 frames. This thread then enters a loop in which it responds to the client's requests until it detects that the client has disconnected. We have defined a very simple protocol for submitting such requests in which the client either requests to alter encoding parameters or stop the video stream. To update the encoder, the client sends the following message:

```
<request action="start">
        <width>#</width>
        <height>#</height>
```

```
        <fps>#</fps>
        <rate>#</rate>
</request>
```

where "width" is the new desired width, "height" is the new desired height, "fps" is the new desired frame rate, and "rate" is the new desired video bitrate. To stop the encoder, the following message is sent with the request action as stop:

```
<request action="stop" />
```

Upon disconnecting, the thread processing the client's requests will shut down the encoder, stop the video stream, and exit.

6 Experimental Results

To validate the cognitive video scaling solution, we have developed a test bed as shown in Figure 9.


Fig. 9 Camera System Test Bed

We created a point to point connection between the client and server by setting up the client machine as a DHCP server and connecting it directly to the embedded video server, allocating it an IP address on an arbitrary subnet. We are using dummynet , a widely used network emulator, to emulate the behavior of internet in the lab environment. With dummynet, one can control the traffic over a specific channel by limiting bandwidth, inserting packet losses, inserting delay, etc. In order to simulate bandwidth change detection the client simply reads from a file that contains channel bandwidth information. In our test set up we developed a test application which simultaneously sets the bandwidth of the channel to varying values at certain intervals using dummynet, and writes this bandwidth to a flat file that the client can read from. With this, we are able to implement some of the conditions of a real network and be aware of the bandwidth in the client application.

The first experiment tests the ability of the system to control the bandwidth of the video by altering the spatial resolution. Channel bandwidth is kept constant, as well as frame rate which is kept at 30 frames per second.

The spatial resolution is changed from 800x480 to 480x320 to 320x240. We used Wireshark to capture and display the bandwidth of the video, obtaining results in Figure 10.


Fig. 10 Spatial Resolution Experiment

As the spatial resolution is changed, the bandwidth of the video changes accordingly. One can easily deduce that this change is directly proportional to the change in resolution. For example, when the spatial resolution changes from 800x480 to 480x320, the total pixel ratio and the ratio of video bandwidth are equivalent:

$$\frac{480 * 320}{800 * 480} = \frac{76800}{384000} = 0.4$$

$$\frac{537.6 Kbps}{1344 Kbps} = 0.4$$

These results indicate that we have successfully demonstrated control over the video bandwidth by altering the spatial resolution of the video.

In our next test, we showed that we can control video bitrate by altering the amount of compression (number of bits per pixel), resulting in a change in video quality. We tested the systems response to changes in channel bandwidth by altering the bandwidth from 1600Kbps to 800Kbps to 400Kbps. Spatial resolution was kept constant at 800x400 and temporal resolution was kept constant at 30 fps. The resulting changes in video bitrate, as well as playback quality, are depicted in Figure 11.
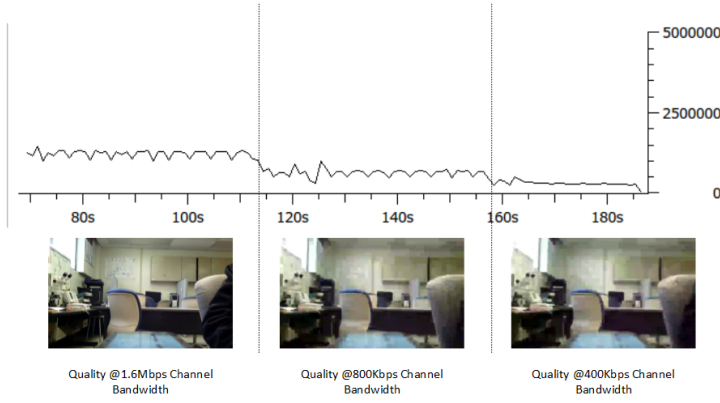
Fig. 11 Quality Resolution Experiment

As can be seen from Figure 15, the system successfully and immediately responds to changes in channel bandwidth by reducing the quality resolution of the video. In addition, the reduction in video bitrate is directly proportional to the reduction in available bandwidth.

Finally, we tested the systems ability to determine the users preferences and scale the video appropriately when channel bandwidth changes. For the purposes of this study, we provided the learning mechanism with an arbitrary training set with the intent to prove the systems ability to properly learn and make predictions. By using a predefined training set we know in advance what classes should be predicted, allowing us to validate the accuracy of the cognitive mechanism by comparing the experimental predicted values with the expected values. The training set used is given in Table 1 and the class distribution for this training set is graphed in Figure 12.


Fig. 12 Class Distribution

The "Content Type" feature is used to classify different scenes that are transmitted in real world scenarios. For example, a content type of 0 may be a video conferencing application with a "talking head" based scene, while a content type of 2 may be a medical based scene that requires extremely high

Table 1 Training Set

Channel Bandwidth (Kbps)	Content Type	Class Label
1200	0	0
1300	0	0
800	0	2
500	0	3
900	0	2
300	0	3
1100	0	0
1300	0	0
1500	0	0
850	0	2
500	1	3
1200	1	0
600	1	1
1300	1	0
300	1	3
400	1	3
800	1	0
600	1	1
700	1	1
1200	2	0
500	2	3
1200	2	0
2000	2	0
1300	2	2
300	2	3
900	2	3
1400	2	2
1900	2	0
700	2	3
1200	2	2

bandwidth and video quality. In this experiment, 3 different content types are used, giving the approximated expected prediction values in Table 2.

Table 2 Expected Prediction Values

Channel Bandwidth (Kbps)	Content Type	Expected Class
0-500	0	3
500-1000	0	2
1000+	0	0
0-500	1	3
500-800	1	0
800+	1	0
0-900	2	3
900-1700	2	2
1700+	2	0

We trained the support vector machines and tested the learning mechanism by setting the bandwidth to 1600Kbps, 800Kbps, and 400Kbps while viewing video streams with content types 0, 1, and 2. For equation (6) we selected the

K parameter to be 0.85 and for equation (8) we selected the Q_{max} parameter to be 3.5. In all the experiments the output bitrates are selected from equation (9). The resulting bandwidth changes, encoding parameter changes, and class predictions are given in Figures 13 - 15.
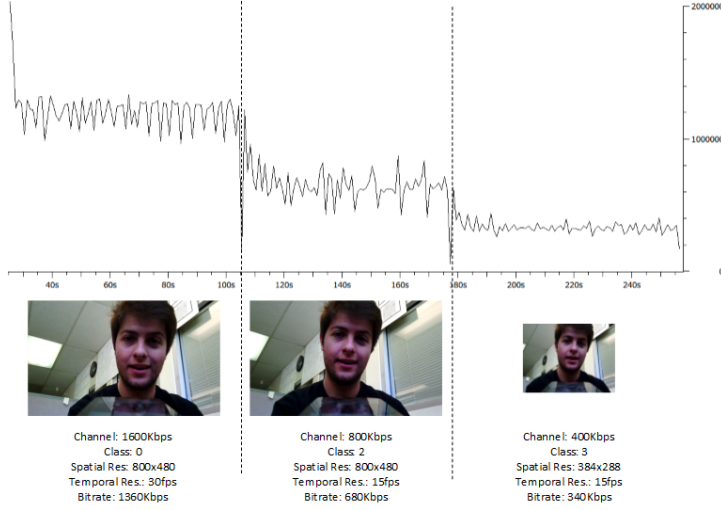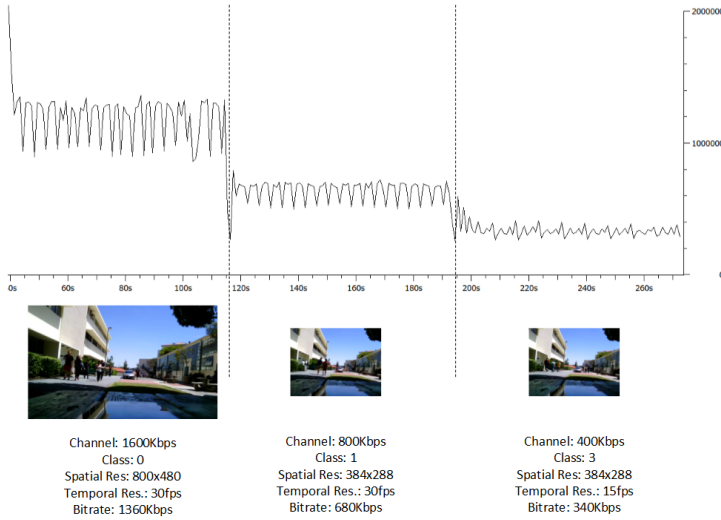

Fig. 13 Predictions With Content Type 0

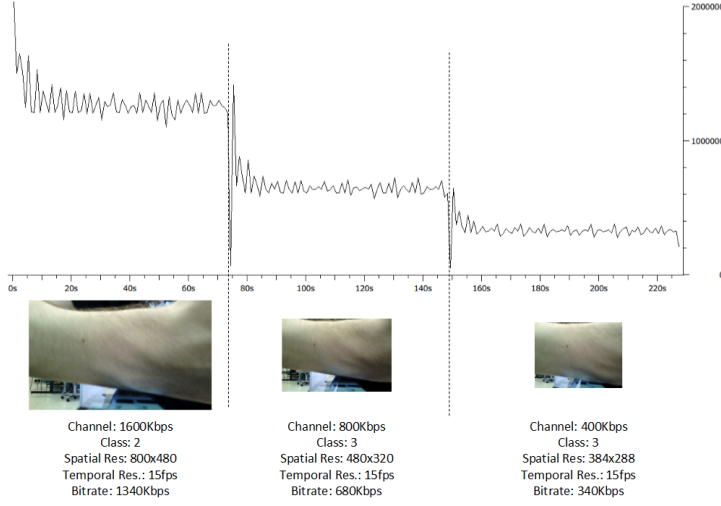**Fig. 14** Predictions With Content Type 1

Fig. 15 Predictions With Content Type 2

In all cases, the video bandwidth was adapted when the channel bandwidth changed. In addition, the video bandwidth was consistently kept at 85% of the channel bandwidth as a result of our selection of K. In order to validate the accuracy of the predictions, Table 3 compares the expected class values and the predicted class values at time instances T_0, T_1, and T_2, where T_0 is the instant when the channel bandwidth changes to 1600Kbps.

Table 3 Expected vs. Predicted Classes

Segment	Content Type	Expected Class	Predicted Class
	0	0	0
T_0	1	0	0
	2	3	3
	0	2	2
T_1	1	1	1
	2	3	3
	0	3	3
T_2	1	3	3
	2	3	3

As the table indicates, the support vector machines predicted the correct class with 100% accuracy. In addition, the resulting changes to the encoding parameters followed the changes defined in Figure 7. When class 0 was predicted, the spatial and temporal resolutions were kept high at the cost of fewer bits per pixel. When class 1 was predicted, the temporal resolution was kept high and the spatial resolution was reduced. The prediction of class 2 resulted in a loss in temporal resolution with the spatial resolution being kept high.

Finally, when class 3 was predicted, the spatial and temporal resolutions were reduced, resulting in greater quality with more bits per pixel.

7 Conclusion

As the volume of internet traffic related to video streaming increases, the importance of having exceptional bitrate adaptation schemes grows. In addition, these schemes should be aware of not only the channel bandwidth, but the surrounding context of the video being streamed. This context encapsulates the content type of the video, and can extend into other dimensions such as amount of motion, geospatial location, and more. We have presented a solution that takes rate adaptation beyond simply changing the quality of the video when the channel bandwidth becomes limited. Our system determines the users preferences about video quality, taking into account if the user prefers a drop in temporal or spatial resolution versus quality resolution. We have demonstrated the ability to adapt to channel bandwidth changes by altering these resolutions. In addition, we have shown that by using support vector machines, we can learn the users preferences and successfully adapt the video bandwidth in line with these preferences. Such a system is a viable solution to the changing atmosphere of video providers, contexts, and the increasing and diverse consumer base.

Acknowledgements

References

1. Http live streaming overview. Apple Inc. Available from: `https://developer.apple.com/library/ios/documentation/NetworkingInternet/Conceptual/StreamingMediaGuide/StreamingMediaGuide.pdf`
2. Adhikari, V.K., Guo, Y., Hao, F., Varvello, M., Hilt, V., Steiner, M., Zhang, Z.L.: Unreeling netflix: Understanding and improving multi-cdn movie delivery. In: 2012 Proceedings IEEE INFOCOM, pp. 1620–1628. IEEE (2012)
3. Chang, C.C., Lin, C.J.: Libsvm: a library for support vector machines. ACM Transactions on Intelligent Systems and Technology (TIST) **2**(3), 27 (2011)
4. De Cicco, L., Mascolo, S., Palmisano, V.: Feedback control for adaptive live video streaming. In: Proceedings of the second annual ACM conference on Multimedia systems, pp. 145–156. ACM (2011)
5. ITU-T RECOMMENDATION, H.: 264 advanced video coding for generic audiovisual services. ISO/IEC **14496** (2003)
6. Lohmar, T., Einarsson, T., Frojdh, P., Gabin, F., Kampmann, M.: Dynamic adaptive http streaming of live content. In: 2011 IEEE International Symposium on a World of Wireless, Mobile and Multimedia Networks (WoWMoM), pp. 1–8. IEEE (2011)
7. Ohm, J.R.: Advances in scalable video coding. Proceedings of the IEEE **93**(1), 42–56 (2005)
8. Oyman, O., Singh, S.: Quality of experience for http adaptive streaming services. IEEE Communications Magazine **50**(4), 20–27 (2012)

9. Ozturk, Y., Kulkarni, M.: Dichirp: direct injection bandwidth estimation. International Journal of Network Management **18**(5), 377–394 (2008). DOI 10.1002/nem.674. URL `http://dx.doi.org/10.1002/nem.674`
10. Pan, C.H., Lee, I.H., Huang, S.C., Lian, C.J., Chen, L.G.: A quality-of-experience video adaptor for serving scalable video applications. Consumer Electronics, IEEE Transactions on **53**(3), 1130–1137 (2007)
11. Schulzrinne, H.: Real time streaming protocol (rtsp) (1998)
12. Schwarz, H., Marpe, D., Wiegand, T.: Overview of the scalable video coding extension of the h. 264/avc standard. IEEE Transactions on Circuits and Systems for Video Technology **17**(9), 1103–1120 (2007)
13. Stockhammer, T., Hannuksela, M.M., Wiegand, T.: H. 264/avc in wireless environments. IEEE Transactions on Circuits and Systems for Video Technology **13**(7), 657–673 (2003)
14. Unane, I., Urteaga, I., Husemann, R., Del Ser, J., Roseler, V., Rodriguez, A., Sanchez, P.: A tutorial on h.264/svc scalable video coding and its tradeoff between quality, coding, efficiency, and performance. In: J.D.S. Lorente (ed.) Recent Advances on Video Coding. InTech (2011). Available from: `http://www.intechopen.com/books/recent-advances-on-video-coding/`
15. Wenger, S.: H. 264/avc over ip. IEEE Transactions on Circuits and Systems for Video Technology **13**(7), 645–656 (2003)
16. Wenger, S., Stockhammer, T.: Rtp payload format for h. 264 video (2005)
17. Wiegand, T., Sullivan, G.J., Bjontegaard, G., Luthra, A.: Overview of the h. 264/avc video coding standard. IEEE Transactions on Circuits and Systems for Video Technology **13**(7), 560–576 (2003)
18. Zambelli, A.: Iis smooth streaming technical overview. Microsoft Corporation **3** (2009)